Problem Statement:

Extend your Lab4 (RV64I simulator) code to support cache simulation to identify which accesses would cause a cache hit and which accesses would cause a cache miss when an assembly program is executing. To support cache simulation within your simulator, the following new commands starting with cache_sim should be added to the Lab4 submission:

1. cache_sim enable config_file - Enables cache simulation for D-cache within the simulator. Uses parameters from config_file for cache settings (see format below).
2. cache_sim disable - Disables cache simulation. This is the default state at the start of the simulator.
   The above 2 commands can not be executed while a file is being executed or loaded.
3. cache_sim status - prints the enable/disable status of the D-cache simulation. If cache simulation is enabled, it also prints the cache configuration values in the format:

   Cache Size: 32168
   Block Size: 16
   Associativity: 8
   Replacement Policy: LRU
   Write Back Policy: WT

4. cache_sim invalidate - invalidates all entries of the D-cache.
5. cache_sim dump myFile.out - writes all the current D-cache entries to the file "myFile.out" in the following format. Notice that only the valid entries have been printed in order of their Set value.

   Set: 0x00, Tag: 0x100, Clean
   Set: 0x01, Tag: 0x123, Clean
   Set: 0x01, Tag: 0x736, Dirty
   Set:0x02, Tag: 0x145, Dirty
   Set:0x10, Tag: 0x321, Clean

6. cache_sim stats - prints the cache statistics for the executing code at the current instance in the format:

   D-cache statistics: Accesses=100, Hit=93, Miss=7, Hit Rate=0.93

If cache simulation is enabled, your program should generate a file named ~~cache.output~~ **filename.output** where filename is the name of the executed file without the extension**,** which should contain cache simulation data in a format as shown:

Example output:
      R: Address: 0x20202, Set: 0x02, Miss, Tag: 0x202, Clean
      W: Address: 0x10306, Set: 0x06, Hit, Tag: 0x103, Dirty
      R: Address: 0x20511, Set: 0x11, Miss, Tag: 0x205, Clean

Also, when invoking the simulator's "run" command, you should print the cache statistics at the end.

If you run multiple files in a single simulator instance, say file1 and file 2, the output file should be generated independently for each executed file.
The output file should be named file1.output and file2.output.

The program should allow customization for various cache parameters as follows, which are provided as a config file in a specific format:
    a. Total size of the cache (specified in Bytes)
    b. Size of each cache block (specified in Bytes)
    c. Associativity (1: Direct mapped, 0: Fully associative, any other number: set associative) - up to 16
    d. Replacement policy (FIFO, LRU, RANDOM)
    e. Write policy (WB: WriteBack, WT: WriteThrough). Assume No-Allocate for WT and Allocate for WB.

The cache configuration is provided as an input file (e.g., cache.config) in the following format:
      SIZE_OF_CACHE (number)
      BLOCK_SIZE (number)
      ASSOCIATIVITY (number)
      REPLACEMENT_POLICY (FIFO or LRU or RANDOM)
      WRITEBACK_POLICY (WB or WT)
Example config file:
      32168
      16
      8
      LRU
      WT

For simplicity (if it helps), assume the following bounds and steps:
    a. Cache size: max. 1 MB, always a power of 2
    b. Block size: max. 64 Bytes, always a power of 2
    c. Associativity: max. 16, always a power of 2
    d. Input address: always 20 bits long. Assume it to be a 32-bit integer with the first 12 bits set to 0. So, only 20-bits are valid values.

You could approach the problem in multiple sub-parts, as mentioned below, with partial weightages for each part.

1. **Part-1:** Only support read access modeling for a direct-mapped cache
2. **Part-2:** Part-1 + caches with associativity + FIFO, LRU and RANDOM replacement policies
3. **Part-3:** Part-2 + write access modeling

Submission instructions:
The assignment is to be done in the same project group as the Lab4. Submit your code and a Makefile, which will compile your code and generate an executable named riscv_sim. Prepare a short report on your coding approach and what all you did to test your code to be correct. You should clearly indicate which parts (part-1, 2, 3) are supported by your code.
ONLY ONE MEMBER OF A GROUP SHOULD MAKE THE SUBMISSION.
Submit a zip file (Lab7_CSYYBTECHZZZZZ_CSYYBTECHXXXXX.zip) containing the following:
1. Source files of your code
2. Makefile - A makefile that would compile your code and generate an executable named riscv_sim.
3. README - Containing information about various files in the directory, usage instructions, etc.
4. Test Cases - if you tried out specific input files.
5. report.pdf - a short report on your implementation/coding approach and what all you did to test your code to be correct.

Demo: We may later conduct a demo where you will show the functioning of your code to the TAs and accordingly be assigned marks.

Examples:
1. **Input1.s:**

```
.data
.dword 10, 20, 30, 40, 50
.text
lui x3, 0x10
ld x4, 0(x3)
ld x4, 8(x3)
ld x4, 16(x3)
ld x4, 24(x3)
ld x4, 32(x3)
```

config.txt:

```
256
16
1
LRU
```

Command Line Interface:

```
cache_sim enable config.txt
load input.s
run
Executed ... PC=...
Executed ... PC=...
Executed ... PC=...
Executed ... PC=...
Executed ... PC=...
Executed ... PC=...
D-cache statistics: Accesses=5, Hit=2, Miss=3, Hit Rate=0.40
cache_sim dump filename.ext
```

input1.output:

```
R: Address: 0x10000, Set: 0x0, Miss, Tag: 0x100, Clean
R: Address: 0x10008, Set: 0x0, Hit, Tag: 0x100, Clean
R: Address: 0x10010, Set: 0x1, Miss, Tag: 0x100, Clean
R: Address: 0x10018, Set: 0x1, Hit, Tag: 0x100, Clean
R: Address: 0x10020, Set: 0x2, Miss, Tag: 0x100, Clean
```

filename.ext:

```
Set: 0x0, Tag: 0x100, Clean
Set: 0x1, Tag: 0x100, Clean
Set: 0x2, Tag: 0x100, Clean
```

2. **Input2.s**

```
.data
.dword 20, 30, 40, 50, 60
.text
lui x3, 0x10
sd x0, 0(x3)
ld x4, 8(x3)
ld x4, 0(x3)
ld x4, 16(x3)
ld x4, 24(x3)
```

config.txt:

```
1024
16
1
FIFO
WT
```

Command Line Interface:

```
cache_sim enable config.txt
load input.s
run
Executed ... PC=...
Executed ... PC=...
Executed ... PC=...
Executed ... PC=...
Executed ... PC=...
Executed ... PC=...
D-cache statistics: Accesses=5, Hit=2, Miss=3, Hit Rate=0.40
```

input2.output:

```
W: Address: 0x10000, Set: 0x0, Miss, Tag: 0x40, Clean
R: Address: 0x10008, Set: 0x0, Miss, Tag: 0x40, Clean
R: Address: 0x10000, Set: 0x0, Hit, Tag: 0x40, Clean
R: Address: 0x10010, Set: 0x1, Miss, Tag: 0x40, Clean
R: Address: 0x10018, Set: 0x1, Hit, Tag: 0x40, Clean
```

Clarifications and Questions:

1. What should be written in the Clean/Dirty column if the policy is WT?
   If the policy is WT, all cache entries can be assigned Clean.
2. You can assume that no read/write will span more than 1 cache block. You can emit an error message for an erroneous input.
3. ~~The entries of the cache have to be invalidated along with the memory at the end of execution of a file.~~ The cache is cleared along with the registers whenever a new .s file is loaded in the simulator.
   (https://moodle.cse.iith.ac.in/moodle/mod/forum/discuss.php?d=294)
4. You also have to simulate the memory in accordance with the cache. For example, if we have a write-back policy, any dirty memory block would have an older value in memory.
5. In the case of write through no allocate policy, any write to a memory block will not go through or be stored in the cache.